

## Implementation Notes

**What this document is:** a complete record of what was built, why each decision was made, and what each change was worth — written for someone with **no background in search or recommender systems**.

Every technical term is explained in plain language, with a real example from this catalog, before it is used. If you hit a word you don't recognise, it is either defined in §2 Vocabulary or in the Glossary at the end.

Each stage section ends with “**Ideas for this stage**” — enhancements specific to that piece of the pipeline. Enhancements that cut across several stages, or that are about how the stages are *orchestrated*, are collected in §10 at the end.

**Result:** the supplied starter agent scores **0.1067**. This implementation scores **0.8592** — roughly 8x — using no AI model, no network connection, and nothing outside the Python standard library.

---

### Table of contents

1. What the agent has to do
2. Vocabulary
3. How the score is calculated
4. Why the supplied baseline scored 0.1067
5. Stage-by-stage changes
  - S0 — Experiment harness
  - S1 — Catalog index
  - S2 — Intent router
  - S3 — Dialog state
  - S4 — Clarification policy
  - S5 — Retrieval
  - S6 — Reranking
  - S7 — Recommendation timing
  - S9 — Robustness and tests
6. What was tried and rejected
7. Results
8. Why upstream Amazon data was ruled out
9. Files at a glance
10. Cross-cutting and orchestration ideas
11. Glossary

---

### 1. What the agent has to do

There is a **catalog** of 50,000 clothing, shoe and jewellery products (`data/catalog.jsonl`). One of them is the **target** — a specific item a

specific customer actually bought. The agent is never told which one.

A **session** is one conversation. It runs for at most **10 turns**. On each turn:

- the customer sends a message,
- the agent replies with a sentence, optionally *a question*, and optionally *a list of up to 10 product IDs*.

The session ends the moment the target appears anywhere in that list — or after turn 10, which counts as a failure.

Here is a real session from the public set, produced by the finished agent:

session public\_0008

TARGET: Hanes Womens Wireless Bra, Full-Coverage Pullover Stretch-Knit Bra

```
turn 1  customer> I'm looking for Bras Everyday Bras. A key requirement is: nylon.
        agent   > To narrow this down: is there anything else that matters
                  for this one?                                [asks: other]  no list yet

turn 2  customer> For that, what matters is: 96% Nylon, 4% Spandex; Pull-On closure.
        agent   > (same question)                             [asks: other]  no list yet

turn 3  customer> For that, what matters is: Hand Wash Only.
        agent   > (same question)                             [asks: other]  10 products
                  -> TARGET at position 1.  HIT on turn 3.
```

*(That is a real run, reproducible with the agent as it stands. Not every session looks like this — in some, the agent becomes confident enough to show a list on turn 2; §57 explains when and why.)*

Two things in that transcript are the whole story of this project, and both are explained in detail later:

1. The customer **only volunteered information because the agent asked a question**. Had the agent not asked, turns 2 and 3 would have contained nothing useful.
2. The agent **deliberately showed no products on turns 1 and 2**, even though it could have. §3 explains why holding back is worth more than it costs.

---

## 2. Vocabulary

Read this once; everything later depends on it.

**The data**

| Term                  | Meaning                                                                                                                                                                                                                                                                                     |
|-----------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Catalog</b>        | The 50,000 products. One JSON object per line in <code>data/catalog.jsonl</code> .                                                                                                                                                                                                          |
| <b>parent_asin</b>    | Amazon's product ID, e.g. B07KCFS4VC. The <i>only</i> thing that is scored — a recommendation is right only if this string matches the target exactly. Close is not counted.                                                                                                                |
| <b>Product fields</b> | <b>title</b> , <b>features</b> (bullet points), <b>description</b> , <b>price</b> , <b>categories</b> , <b>details</b> (a dictionary like <code>{"Department": "Womens"}</code> ), <b>average_rating</b> , <b>rating_number</b> , <b>store</b> (the brand).                                 |
| <b>Session</b>        | One conversation with one customer, targeting one product.                                                                                                                                                                                                                                  |
| <b>Turn</b>           | One customer message plus one agent reply. Maximum 10 per session.                                                                                                                                                                                                                          |
| <b>Target</b>         | The hidden product the agent must find.                                                                                                                                                                                                                                                     |
| <b>Scenario</b>       | Sessions come in four flavours: <b>buying</b> (customer states a firm requirement immediately), <b>browsing</b> (customer starts vague), <b>intent override</b> (customer changes their mind partway through), <b>boundary</b> (customer has no opinion on something the agent asks about). |

### The two halves of any search system

This distinction is the single most important idea in the document, and the architecture is built around it.

**Retrieval** = *finding* a shortlist of plausible candidates out of 50,000.

**Ranking** = *ordering* that shortlist so the best one is first.

They are different jobs with opposite requirements:

- Retrieval must be **generous**. If the target isn't in the shortlist, nothing downstream can save you. Retrieval is allowed to include junk.
- Ranking must be **strict**. Its job is to separate the one right answer from 299 near-misses.

A technique can be excellent at one and terrible at the other. §6 documents a case where exactly that happened, and where moving one technique from

retrieval to ranking turned a large loss into a large gain.

### How text search works (BM25)

The agent finds products by matching words. The scoring function used is **BM25**, a standard text-search formula. You don’t need the formula — you need three intuitions:

1. **Rare words count more than common words.** If a customer says “kandinsky”, that word appears in maybe one product in the whole catalog, so it is enormously informative. If they say “imported”, which appears in tens of thousands, it tells you almost nothing. BM25 weighs each word by its rarity across the catalog. (The technical name for this rarity weighting is **IDF**, *inverse document frequency*.)
2. **More matches are better, but with diminishing returns.** A product mentioning “leather” five times isn’t five times more relevant than one mentioning it once.
3. **Which field matched matters.** A word in the product’s title is stronger evidence than the same word buried in a long **description**. This is configurable — see `DEFAULT_WEIGHTS` at `src/index.py:25`.

The search index itself is **SQLite FTS5**, a full-text search engine built into Python’s standard library. That is why this project needs no third-party packages: the search engine is already on your machine.

### Measuring success

| Term                        | Meaning                                                                                                                                                                                                       |
|-----------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Top-k</b>                | The first $k$ items of a ranked list. Here $k = 10$ .                                                                                                                                                         |
| <b>Hit</b>                  | The target appeared in a list the agent showed.                                                                                                                                                               |
| <b>Hit Rate@10</b>          | Fraction of sessions that ended in a hit. <b>0.94</b> = the agent found the target in 94% of sessions.                                                                                                        |
| <b>Rank</b>                 | The target’s position in the list. Position 1 is best.                                                                                                                                                        |
| <b>Reciprocal rank (RR)</b> | $1 \div \text{rank}$ . Rank 1 $\rightarrow 1.0$ . Rank 2 $\rightarrow 0.5$ . Rank 10 $\rightarrow 0.1$ . A miss $\rightarrow 0$ . This rewards being <i>first</i> far more than merely being <i>present</i> . |
| <b>MRR</b>                  | Mean Reciprocal Rank — the average RR across all sessions. <b>0.79</b> roughly means “on average the target lands around position 1.3”.                                                                       |

| Term        | Meaning                                                                                                                  |
|-------------|--------------------------------------------------------------------------------------------------------------------------|
| <b>MTTC</b> | Mean Turns To Conversion — the average turn number on which the hit happened. A miss counts as turn 11. Lower is better. |

#### Terms specific to this project

| Term                   | Meaning                                                                                                                                                                                                                                         |
|------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Facet</b>           | A structured attribute of a product with a limited set of values — material ( <b>leather</b> , <b>cotton</b> ), colour ( <b>black</b> , <b>navy</b> ), price band, category. Contrast with free text. Extracted by <code>src/facets.py</code> . |
| <b>Constraint span</b> | A short phrase the customer said, normalised for matching — e.g. " <b>stainless steel band</b> ". §S6 explains why these are so powerful here.                                                                                                  |
| <b>ask_attribute</b>   | A structured field in the agent's reply naming <i>what</i> it is asking about ( <b>material</b> , <b>color</b> , <b>budget</b> , <b>other</b> , ...). Critically, this — not the English sentence — is what the customer's simulator reads.     |

### 3. How the score is calculated

All of this lives in `evaluator/local_evaluator.py`, which is organizer-owned and **must never be modified**.

#### Step 1 — cleaning the agent's list

`normalize_recommendations` (`evaluator/local_evaluator.py:95`) silently discards:

- IDs that aren't in the catalog,
- duplicates,
- anything after the first 10.

“Silently” matters: a malformed recommendation isn't an error, it just quietly wastes a slot.

## Step 2 — checking for a hit

```
if override_applied and target in ranked:
    best_rank = ranked.index(target) + 1
    hit_turn = turn
    break
```

— evaluator/local\_evaluator.py:252

Note the **break**. **The session stops at the first hit.** This has a consequence that shapes the entire agent design, covered below.

## Step 3 — the four metrics

From `metric_summary` (evaluator/local\_evaluator.py:188) and line 280:

```
HitRate@10 = hits / N
MRR         = average(1 / rank), a miss contributes 0
MTTC        = average(hit turn), a miss counts as turn 11
Efficiency = (11 - MTTC) / 10, clamped to the range 0..1
```

$\text{Score} = 0.50 \times \text{HitRate@10} + 0.30 \times \text{MRR} + 0.20 \times \text{Efficiency}$

For this implementation:

$0.50 \times 0.940 + 0.30 \times 0.7911 + 0.20 \times 0.7595 = 0.8592$

## The five consequences that drive every design decision

- 1. The first list containing the target freezes both rank and turn.** Because of that **break**, you never get a second chance to rank better. If you show a list on turn 2 and the target happens to be sitting at position 9, you have permanently locked in  $\text{RR} = 0.11$  — even though by turn 4 you might have ranked it first. This is the single most important strategic fact in the challenge.
- 2. Wrong recommendations are free.** There is no penalty for showing 10 irrelevant products. The only cost is the turn. So the dangerous failure mode is not *being wrong* — it is *being confident too early*.
- 3. ask\_attribute is the only tap for information.** The simulated customer (`customer_reply`, evaluator/local\_evaluator.py:166) reveals a new constraint **only** when the agent's reply sets `ask_attribute`. If it is `null`, the customer answers “ask me about one specific attribute” and discloses nothing. An agent that never asks never learns anything.
- 4. An exception is an invisible failure.** The evaluator wraps each turn in `try / except Exception` (evaluator/local\_evaluator.py:241) and converts a crash into an empty response. You lose the session and see no error message.

**5. Override sessions cannot convert early.** For intent-override scenarios, `override_applied` stays `False` until the customer’s change of mind arrives on turn 3 or 4. Hits before that are ignored entirely, which puts a hard floor under that scenario’s MTTC.

### What each improvement is actually worth

With 200 sessions, here is the score change from improving a *single* session:

| Change to one session                       | Score change    |
|---------------------------------------------|-----------------|
| Turn a miss into a hit (rank 5, turn 4)     | <b>+0.0035</b>  |
| Move an existing hit from rank 10 to rank 1 | <b>+0.00135</b> |
| Reach the same hit one turn sooner          | <b>+0.0001</b>  |

Read those numbers together and they give a clear instruction:

**Converting a miss is ~2.6x more valuable than fixing a rank, and fixing a rank is ~13x more valuable than the turn it costs to earn.**

Therefore: *spend turns freely to gather information; never rush a list out.*

Almost every design choice in §5 follows from this one table. It is why the agent stays silent for the first two turns.

---

## 4. Why the supplied baseline scored 0.1067

The original `starter/agent.py` was about 100 lines. Its `respond()` did three things: tokenise **the current message**, run one BM25 query, return the top 10. It scored:

|              |               |
|--------------|---------------|
| Hit Rate@10  | 0.125         |
| MRR          | 0.068         |
| MTTC         | 9.81          |
| <b>Score</b> | <b>0.1067</b> |

Three flaws, in order of cost:

**It never asked a question.** `ask_attribute` was hard-coded to `None`. By consequence #3 above, the customer therefore never disclosed anything. The agent’s entire knowledge of the customer was the opening sentence — for all ten turns.

**It was stateless.** Even if the customer *had* said something useful, `respond()` only ever looked at `user_message`, the current turn. Everything said earlier was discarded.

**It never reranked.** Whatever order BM25 returned was the order shown.

The MTTC of 9.81 tells the story: out of 10 possible turns, the average session burned 9.81 of them. The agent was effectively making the same guess ten times in a row, because nothing ever changed between turns.

**The key realisation:** the biggest available win was not a better search engine. It was *asking a question and remembering the answer*. Fixing only that — before touching retrieval or ranking at all — took the score from **0.1067 to 0.7811**.

---

## 5. Stage-by-stage changes

The agent is split into one module per job, under `src/`. `starter/agent.py` remains the file the evaluator imports; it is now a thin adapter that wires the stages together and owns the response format.

Reading order for the code: `starter/agent.py` → `src/state.py` → `src/retrieval.py` → `src/rerank.py`.

---

### S0 — Experiment harness

**File:** `tools/sweep.py`

**What this stage does** It isn't part of the agent. It is the tool used to *measure* the agent — run several configurations and print a comparison table.

**Why it was needed** Two problems made experimentation painful.

*Speed.* Building the search index over 50,000 products takes about 4 seconds, and the original setup rebuilt it for every configuration tested. Comparing four ideas meant four rebuilds, and early runs timed out at two minutes.

*Self-deception.* There are only 200 public sessions, but **800 private sessions decide the final result**. If you tune ten parameters against 200 sessions and keep whatever scores highest, you will reliably pick settings that fit the noise in those particular 200 — and they will not transfer. This is called **overfitting**.

**What changed** **Shared index.** The index is now cached per catalog path (`load_index` in `src/index.py`), so constructing many `Agent` objects builds it once. A full 200-session evaluation dropped from timing out to **~13 seconds**.



**Dev/holdout split.** `split_samples()` divides the 200 sessions into 120 **dev** and 80 **holdout**, stratified so each split keeps the same mix of scenarios. All tuning happens on dev. Holdout is looked at only to confirm a change survives on data it wasn't tuned against. Where dev and holdout disagreed, holdout won — twice, this changed what shipped (see §6).

The split is deterministic (sorted by `sample_id`), so it needs no stored file and never drifts.

```
python3 tools/sweep.py --split dev
python3 tools/sweep.py --split holdout
python3 tools/sweep.py --split all --configs floor,rerank,infogain
```

**Measured effect** No direct score. Every measurement in this document exists because of it.

#### Ideas for this stage

- **Confidence intervals.** With 80 holdout sessions, a difference under about 0.02 is indistinguishable from noise. The harness reports point estimates only. Bootstrap resampling over sessions would print an interval, turning “is this better?” from a judgement call into a measurement.
- **Per-session diffing.** The harness compares aggregates. A mode that lists which *specific* sessions changed between two configurations would make regressions far faster to diagnose — currently that requires an ad-hoc script.
- **Automated parameter search.** Several constants were swept by hand (§S7, §S4). A small random or grid search driven from the config matrix would cover more ground, though it raises overfitting risk and would need the confidence intervals above to be safe.
- **Cache the catalog parse.** `catalog_index()` re-parses 58MB of JSON per harness invocation (~2s). A pickled cache would speed up repeated runs from the shell.

---

#### S1 — Catalog index

**Files:** `src/index.py`, `src/text.py`

**What this stage does** Reads the 50,000-product catalog once and builds the structures every later stage queries: a full-text search index, and a trimmed copy of each product's data.

**Why it was needed** The original code built its index inline inside the agent and kept nothing else, so no other component could ask questions like “what

colour is this product?” — the information was inside the search engine and unreachable.

**What changed** **One pass, three products.** `CatalogIndex._build()` walks the file once and produces:

1. The **FTS5 full-text index** — seven columns (`title`, `categories`, `features`, `details`, `store`, `description`), so field-specific weighting is possible.
2. **Trimmed product records** — `parent_asin`, `title`, `categories`, `brand`, `price`, `ratings`, and a normalised text blob. Deliberately not a second full copy of the 58MB file.
3. **Normalised text** — this one is subtle and it matters. Product text is stored with all punctuation stripped and words joined by single spaces:

"Hide & Drink, 100% Leather, Buckle closure"

becomes

"hide drink 100 leather buckle closure"

The reranker (§S6) checks whether a phrase the customer said appears inside a product’s text. Without matching normalisation on both sides, "100 leather" would never match "100% Leather" and the strongest signal in the system would silently do nothing.

**Shared caching.** `load_index()` returns one index per catalog path, so repeated **Agent** construction is nearly free.

**Word handling** (`src/text.py`). Two lists of words are removed before searching:

- *Stopwords* — conversational filler: “the”, “for”, “looking”, “preference”.
- *Boilerplate* — Amazon catalog filler: “imported”, “closure”, “machine”, “wash”, “dimensions”. These appear in tens of thousands of products, so they add query cost without discriminating.

**Measured effect** Index build **4.13s**, one time at startup. Peak memory **~282MB**. Per-turn latency **16.6ms mean, 19.8ms at the 95th percentile**.

**Ideas for this stage**

- **Tune the field weights.** `DEFAULT_WEIGHTS` at `src/index.py:25` is (0.0, 6.0, 4.0, 2.5, 2.5, 1.5, 1.0) — inherited verbatim from the supplied baseline and **never tuned**. Those seven numbers control how much a `title` match counts versus a `description` match. This is probably the cheapest untapped gain in the repo: the harness exists, the parameter is a single tuple, and nobody has ever checked whether the baseline’s guess was any good.

- **Learn the boilerplate list instead of writing it.** ~~Compute document frequency across the catalog and drop the top 200 terms.~~ **Done, but not the way this originally proposed — the original proposal was wrong.** Ranked by document frequency, **polyester** is 72nd, **cotton** 83rd, **black** 97th, **leather** 111th and **spandex** 141st. Dropping the top 200 would delete every one of them, and those are precisely the constraints the customer discloses, because `intent_card()` inserts a material at position 0 and a colour at position 1 of every card (§S1 explains why the customer quotes catalog copy). Meanwhile **asin** — a genuine member of the hand-written list — sits at rank 10,379 with 0.0% document frequency. Frequency fails in both directions.

What works is *where* a token occurs, not *how* often. Amazon’s structural metadata lives in the **details** dict, and those tokens appear almost nowhere else — **department** 100% of its occurrences, **dimensions** 99.6%, **manufacturer** 99.6%, **inches** 97.7% — while attribute values are spread across title, features and description: **spandex** 1.2%, **cotton** 2.5%, **polyester** 3.3%, **black** 13.4%. The gap between 16% and 96% is empty of real attribute words, so the threshold is read off the catalog’s own distribution rather than tuned against a score. `tools/build_stoplist.py` applies that rule and writes `src/stoplist.py`; it reads `data/catalog.jsonl` and never opens `data/public_set.jsonl`, so it cannot fit the public sessions. The learned list reproduces 14 of the 24 hand-written terms and finds 22 more that nobody thought to write down — every month name and year 2014–2022, harvested from “Date First Available: August 15, 2019”.

**Half the list is not learnable, and `src/text.py` now says so.** **imported**, **machine**, **wash**, **closure** and **care** are care-and-origin phrases in **features/description**. Two statistics were tested and both fail to separate them from real attributes: by document frequency **closure** (38.6% of products) outranks **polyester** (21.8%), and by spread across the 12 largest category buckets **polyester** (CV 0.51) falls between **imported** (0.49) and **wash** (0.60), with **black/white** (0.34/0.31) landing inside the scaffolding band. What separates them is that a shopper does not choose between “imported” and “not imported” — semantics, not frequency. Those ten stay hand-written, with the evidence recorded beside them so the experiment is not retried.

**Measured effect: none.** Public set identical to six decimal places (0.912526 both ways), adversarial set −0.0001, dev and holdout flat. Boilerplate removal strips a median of one token from a ten-token query and `MAX_QUERY_TERMS = 60` never binds — 0 of 200 sessions come close. It ships for auditability and for robustness on a catalog nobody has hand-inspected, not for points. Measurements were taken in one process with everything else held constant; see `tests/test_stoplist.py` for the invariants that keep a future regeneration honest.

- **Persist the index.** 4.1s of startup could become near-zero by writing the FTS5 index to a file instead of `:memory:`. Only worth it if startup time is ever judged.
- **Dead code.** `search_phrases()` (`src/index.py:119`) has had no caller since the phrase route was removed in §6. It should go.

---

## S2 — Intent router

**File:** `src/router.py`

**What this stage does** Classifies the opening message as **buying** (customer knows what they want) or **browsing** (customer is exploring).

**Why it was needed** The problem brief calls for “dual-track routing” as a core pillar: a decided customer and an exploring one deserve different treatment.

**What changed** Classification uses **linguistic cues**, not exact string matching:

```
BUYING_CUES    = "key requirement | must be | i need | it has to | ..."
BROWSING_CUES = "still exploring | just looking | not sure | open to | ..."
```

This was a deliberate choice with a trap avoided. The local evaluator builds its opening lines from fixed templates (`initial_message`, `evaluator/local_evaluator.py:154`), so matching those exact strings would score perfectly here — and could fail completely on the private sessions if the organizer paraphrases them. Cue matching degrades gracefully; template matching fails silently.

Unknown phrasing defaults to **browsing**, because mistaking a vague customer for a decided one commits to constraints they never stated, whereas the reverse costs at most one extra question.

**Measured effect — and an honest scope reduction** The router was first built to change *retrieval*: a wider candidate pool for browsing (400), a narrower one for buying (200). Measured:

|                          | dev    | holdout |
|--------------------------|--------|---------|
| Router affects retrieval | 0.8715 | 0.8373  |
| Router does not          | 0.8715 | 0.8391  |

**Identical on dev, 0.002 worse on holdout.** The reranker (§S6) already resolves both tracks well, so there was nothing left for routing to fix.

Rather than keep a feature that does nothing and describe it as a pillar, the router was **scoped down** to the job it genuinely does: phrasing. A browsing customer hears “To point you in the right direction: ...”; a buying customer hears “To narrow this down: ...”. That is a real product requirement even where it is not a scoring one, and the measurement is recorded in the module docstring so the next reader doesn’t repeat the experiment.

### Ideas for this stage

- **Route the *question*, not the retrieval.** The measurement showed routing retrieval is pointless, but nobody tested routing the clarification policy. A buying customer who has already stated a hard requirement might be better served by a *narrow* confirming question, while a browser needs broad ones. That is a different lever and remains untested.
  - **Detect scenario, not just track.** The router distinguishes two of four scenarios. Detecting *boundary* customers (people who answer “I have no preference”) early would let the agent stop spending questions on someone who won’t answer them — currently that is only learned after a wasted turn, via `dead_attributes` in `src/state.py`.
  - **Confidence-weighted routing.** `classify()` returns a hard label. Returning a confidence, and blending behaviour when uncertain, would avoid an all-or-nothing decision made on a single sentence.
- 

## S3 — Dialog state

**File:** `src/state.py`

**This is the stage that carries the score.** Everything else in this document is a refinement on top of it.

**What this stage does** Remembers the conversation: what the customer has said, when they said it, what they have declined to answer, and whether they changed their mind.

**Why it was needed** The baseline was stateless. Recall consequence #3 from §3: the customer discloses information only in response to a question, and the baseline never asked one. So the baseline was solving a much harder problem than the one posed — identifying one product in 50,000 from a single vague sentence like “*I’m looking for Watches Wrist Watches*”.

Here is what the agent actually receives across a session once it starts asking:

```
turn 1 "I'm looking for Watches Wrist Watches. Stainless Steel Band"
turn 2 "For that, what matters is: Water Resistant; 3 Year Battery."
turn 3 "For that, what matters is: Day / Date Indicator; Stainless Steel Band."
```

Every session hides **exactly four** constraints, disclosed up to two per answered question. Accumulating them across turns instead of discarding them is the difference between describing “a watch” and describing “*this* watch”.

**What changed** `DialogState` stores every message with **provenance** — the turn it arrived on and a weight. It exposes two views of the same history:

- `full_text()` — everything the customer has said. Maximum recall.
- `focused_text()` — only the currently authoritative turns. Maximum precision.

Until the customer changes their mind, these are identical.

**Intent override — a deliberate deviation from the brief.** The problem statement calls for “slot erasure and rewriting”: when the customer reverses a preference, delete the old one. We **down-weight instead of erasing**, to `PRE_OVERRIDE_WEIGHT = 0.35` (`src/state.py:38`).

The reason is specific to how this evaluator constructs an override. Looking at `behavior_for()` (`evaluator/local_evaluator.py:78-86`), the “discarded” preference is *still drawn from the target product’s own metadata*. Erasing it therefore throws away accurate information about the very item we are trying to find, and costs score. Down-weighting keeps the evidence usable while letting the post-override turns dominate, and `focused_text()` gives retrieval a clean post-override view in case the private simulator uses genuine decoys instead.

This deviation is flagged here, in `README.md`, and in the module docstring, because a reader comparing the code to the brief will otherwise think it is a bug.

**Two other signals the state tracks**, both read by the clarification policy:

- **dead\_attributes** — when the customer says “I don’t have a preference for material”, that attribute is recorded as dead so the agent never asks again. Re-asking a declined question wastes a turn and reads as not listening.
- **productive\_turns** — whether an answer actually disclosed something new. This is how the policy tells the difference between “the customer is still telling me things” and “I have exhausted this line of questioning”.

**Measured effect** **0.1067 → 0.7811**. Roughly 90% of the total improvement in this project comes from this stage plus the decision to ask a question at all.

**Ideas for this stage**

- **Tune the override weight.** `PRE_OVERRIDE_WEIGHT = 0.35` was chosen by judgement and **never swept**. Intent override is the weakest scenario in the final results (0.867 hit rate against 0.94 overall, MTTC 4.60 against 3.41), so this constant sits directly on the largest remaining weakness. Sweeping 0.0 (true erasure) through 1.0 (no override handling) via the harness is a small experiment with a real chance of paying.

- **Per-constraint provenance, not per-turn.** Weight currently applies to a whole utterance. A turn that contains both a reversal *and* new information gets uniformly discounted. Tracking weight per extracted constraint span would be finer-grained and better matches what “slot” rewriting means in the brief.
- **Detect *partial* overrides.** `OVERRIDE_CUES` treats any reversal as global. A customer saying “actually, not in black” is revising one facet, not the whole request. Pairing the reversal cue with the facet it mentions would let the agent revise a single slot.
- **Contradiction detection.** Nothing currently notices if a customer states two incompatible constraints. Detecting that and asking which one wins is both a real product behaviour and a natural clarification trigger.

---

## S4 — Clarification policy

File: `src/policy.py`

**What this stage does** Decides **what to ask about next** — the value placed in `ask_attribute`. Legal values are `category`, `material`, `color`, `size`, `style`, `brand`, `budget`, `feature`, `use_case`, `other`.

**Why it was needed** By consequence #3 in §3, this field is the agent’s only means of learning anything. The baseline left it `null` permanently.

**What changed — two implementations** `FixedPolicy (src/policy.py:40)` — **what ships.** Asks the broadest available question every turn, falling through a ladder (`feature`, `use_case`, `style`, ...) as attributes are declined. Simple and, as measured below, effective.

`InfoGainPolicy (src/policy.py:72)` — **the more interesting design.** Chooses the question that would most reduce uncertainty about which candidate is correct.

The intuition, using colour. Suppose the shortlist holds 300 candidates: 150 black, 150 white. Asking “what colour?” eliminates about half whatever the answer is — a genuinely useful question. Now suppose 299 are black and 1 is white. The same question almost certainly returns “black” and eliminates almost nothing. **The value of a question is how evenly it splits what you’re unsure about.** The formal name for that measure is **entropy**.

The policy scores each attribute as:

`value(attribute) = coverage x gain_ratio x answerability`

Each term earns its place, and two of them exist because the naive version failed:

- **coverage** — the fraction of candidates for which the attribute is even resolvable. This is what stops the agent asking about budget: **78.9% of this catalog has price: null**, so a budget question is unanswerable most of the time.
- **gain\_ratio** — entropy divided by its maximum possible value. *This was a bug fix.* Raw entropy is biased toward attributes with many distinct values. Measured on a real session, brand scored **6.1 bits** against colour’s **2.3** purely because the catalog holds thousands of distinct stores, so the agent opened every conversation by asking which brand the customer wanted. Normalising by `log2(number of distinct values)` removes the bias. This is a classic trap in information-gain methods and is worth knowing about generally.
- **answerability** — the probability a *shopper* can answer. Even after the gain-ratio fix, brand still won: every product has a brand, so coverage was 100% and the split was near-perfect. But few people browsing for a belt can name the brand they want. This term encodes shopper behaviour, not catalog structure — and the customer’s own `preference_tags` nudge it, which is the one place the user profile does real work.

Broad questions (“anything else that matters?”) are scored on the same scale rather than special-cased: they can be answered as long as something remains undisclosed, and they return whichever fact the customer thinks is most important, so their value is the average gain across resolvable attributes scaled by how many facts a broad answer tends to surface.

#### Measured effect — and why the simpler one ships

|                | dev           | holdout       |
|----------------|---------------|---------------|
| FixedPolicy    | <b>0.8738</b> | <b>0.8374</b> |
| InfoGainPolicy | 0.8369        | 0.8141        |

`InfoGainPolicy` actually finds the target *slightly more often* on dev (hit rate 0.958 vs 0.950) — but ranks it considerably worse (MRR 0.703 vs 0.814). The cause is understood: specific questions surface fewer facts per turn than broad ones, and thinner evidence ranks worse even when it still retrieves.

The decision rule was fixed **before** the measurement: ship whatever wins on data that was not tuned against. `FixedPolicy` ships. `InfoGainPolicy` stays in the repo, fully documented and selectable:

```
Agent(catalog, AgentConfig(policy=InfoGainPolicy(agent.facets)))
```

#### Ideas for this stage

- **Model expected *yield*, not just split quality.** This is the identified fix for `InfoGainPolicy` and follows directly from the measurement above.



It currently estimates *how much a known answer would narrow the field* but not *how many facts the answer will contain*. Broad questions win here precisely because they return two constraints instead of a fraction of one — and the policy cannot see that. Estimating yield per attribute from observed disclosure sizes, rather than the current fixed `broad_yield` constant, is the single change most likely to make the principled policy overtake the simple one.

- **Ask about the top candidates, not the whole pool.** Entropy is computed over 150 candidates weighted by score. Concentrating on the top ~20 — the ones actually in contention — would ask questions that separate the leaders rather than the field.
- **Two-step lookahead.** The policy is greedy, picking the best single question. Some pairs of questions are worth more together than either is alone.
- **Question phrasing from the candidates.** `QUESTION_TEXT` is a fixed table. Generating “are you after leather or canvas?” from the actual values present in the shortlist would be more natural and more informative — and needs no model, just the facet distribution already computed.

---

## S5 — Retrieval

File: `src/retrieval.py`

**What this stage does** Produces the shortlist of ~300 plausible candidates from 50,000, given everything the customer has said.

**Why it was needed** The baseline ran a single query over the current message only. With state (§S3) the query can now be built from the whole conversation.

**What changed** Two **routes** — independent searches — whose results are combined:

1. **Terms route** — all the customer’s words, over the whole conversation. High recall.
2. **Focused route** — the same, but over post-override turns only. Runs only when the customer has changed their mind, and gives the reversal a clean channel.

**Combining with Reciprocal Rank Fusion (RRF).** Two searches return two ranked lists with scores on incompatible scales — one might score 34.4, the other 18.7, and those numbers mean different things. RRF sidesteps calibration by **ignoring the scores and using only the positions**:

```
combined_score(product) = sum over routes of weight / (60 + position)
```

A product at position 1 contributes 1/61; at position 10, 1/70. Appearing high in both lists beats appearing very high in one. Because it needs no calibration,

adding a route later cannot destabilise the existing ones — and if a route returns nothing, it simply contributes nothing.

**Measured effect** The terms route is **already excellent at its job**. Measured over 80 sampled sessions with all four constraints disclosed:

- target present in the shortlist: **80 / 80**
- median position of the target: **1**
- target within the top 10: **69 / 80**

That result set the direction for everything after it. **Retrieval was not the bottleneck — ranking was.** 80/80 of the targets were being found; 11 of them were simply ordered badly. So effort moved to §S6, and a planned dense-vector retrieval route was cancelled as unnecessary.

#### Ideas for this stage

- **Tune weight\_focused.** Set to 0.8 by judgement (`src/retrieval.py:36`) and never swept. Like `PRE_OVERRIDE_WEIGHT`, it sits directly on the weakest scenario.
- **Weight words by recency.** Every word in the conversation currently counts equally, whether said on turn 1 or turn 5. Later disclosures are more specific and arguably deserve more weight — a cheap experiment on top of the existing `Utterance.weight` machinery.
- **A category route.** The opening message names a coarse category (“Jewelry Necklaces”) that currently contributes only as loose words. A dedicated route restricted to that branch of the category tree would suppress cross-category noise. (See also §S6, where the same idea may work better as a ranking signal.)
- **Dense retrieval, if recall ever becomes the bottleneck.** A local sentence-embedding route would catch paraphrases that share no words. It was cut because recall is 80/80, and it would compromise the no-network guarantee — but if the private sessions paraphrase heavily, recall could fall and this becomes relevant again.

---

## S6 — Reranking

**File:** `src/rerank.py`

**The second-largest win in the project**, and the clearest illustration of the retrieval-versus-ranking distinction from §2.

**What this stage does** Takes the ~300 candidates and reorders them so the best is first.

**Why it was needed** From §5: the target was already in the shortlist 80/80 times at median position 1, but only reached the top 10 in 69 of 80. The gap was pure ordering. And ordering is worth a lot — MRR carries 30% of the score, and by consequence #1 in §3 you only ever get one chance at it.

**What changed The insight.** The constraints the customer discloses are not paraphrases. They are **copied verbatim from the target product’s own metadata** — `intent_card()` (`evaluator/local_evaluator.py:52`) builds them directly out of the product’s **features** and **details** fields. So when a customer says “stainless steel band”, those exact words are sitting in the target product’s text.

That makes an unusually strong signal available: check whether each candidate’s text *literally contains* the phrases the customer used.

customer said: "stainless steel band", "day date indicator"

candidate A text: "...analog watch, stainless steel band, day date indicator..." -> 2 matches

candidate B text: "...leather strap chronograph..." -> 0 matches

The final score for each candidate combines three signals:

```
score = span_coverage          (dominant - verbatim phrase matches)
       + normalised_retrieval_score (BM25's opinion, scaled to 0..1)
       + 0.02 x popularity      (tie-break only)
```

Longer phrases count slightly more (`length_bonus`), because a five-word match is rarer than a two-word one.

**Popularity is deliberately near-zero.** `average_rating` and `rating_number` are held at weight 0.02 — a tie-break, nothing more. The target is *one specific person’s purchase*, not a bestseller, so a strong popularity prior would systematically drag the ranking toward famous products and away from the answer. This is a general lesson: popularity is a good prior for “what will most people like” and a bad one for “which item did this person buy”.

**Why this lives in ranking and not retrieval.** Verbatim matching was *first* built as a third retrieval route, using FTS5 phrase queries. It scored **0.6859 against the floor’s 0.7799 — a large loss**. Investigation showed why: as a search route it found the target in only **47 of 80** sessions, against the terms route’s 80/80. Fused into retrieval, it injected 33 sessions’ worth of noise ahead of good results.

The same evidence, applied to a shortlist the terms route had already filled correctly, is pure gain. **Precision-oriented signals belong in ranking; recall-oriented ones belong in retrieval.** This is the concrete example promised in §2.

**Measured effect** 0.7799 → 0.8543. The single largest gain after dialog state.

#### Ideas for this stage

- **Facet agreement — the clearest unused signal in the repo.** FacetStore is constructed at `starter/agent.py:75`, and `src/facets.py` extracts material, colour, size, brand, price band and category for every product — but **no ranking signal reads any of it**. The shipped FixedPolicy doesn't use facets either, so the extractor currently does no work in the shipped configuration. Adding “candidate's material matches what the customer said” to the score is a small change against code that already exists, and it targets MRR, which carries 30%.
- **Category anchoring.** A traced failure makes the case: for a target belt, the customer had disclosed only "buckle closure" and "100 leather" — both generic — and the target sat at **rank 15** behind a dozen other leather belts. Boosting candidates whose category matches the opening message's stated category is the most direct attack on that failure mode.
- **Learn the signal weights.** `span_weight`, `retrieval_weight`, `popularity_weight` and `length_bonus` were all set by hand. With ~176 known-correct sessions to learn from, even simple logistic regression over these features would likely beat hand-tuning — and stays fully offline.
- **A negative-evidence signal.** Nothing currently penalises a candidate that *contradicts* a stated constraint. A customer who said “leather” and a candidate whose material is explicitly canvas should be pushed down, not merely left unrewarded.
- **What was already tried and rejected here:** weighting phrases by how rare they are within the shortlist. It moved dev by 0.0002 and holdout not at all — a pool retrieved by those same words has little rarity spread left to exploit. The code was deleted rather than kept as a dead option; see §6.

---

#### S7 — Recommendation timing

**File:** `starter/agent.py` (`_shortlist`, `_confident`)

**What this stage does** Decides whether to show a list this turn, or stay quiet and ask another question.

**Why it was needed** This stage exists entirely because of consequence #1 in §3: **the session ends at the first hit, freezing the rank forever**. Showing an uncertain list early risks catching the target at position 9 and locking in  $RR = 0.11$ , when waiting two more turns might have ranked it first.

The marginal-value table in §3 quantifies the trade: a rank improvement is worth ~13x the turn it costs. So the agent should be patient.

**What changed** **Hold until turn 3.** No recommendations on turns 1-2, when the customer has disclosed little. Measured on holdout:

| First recommendation on | Score         |
|-------------------------|---------------|
| turn 2                  | 0.8088        |
| <b>turn 3</b>           | <b>0.8349</b> |
| turn 4                  | 0.8196        |

Turn 3 wins on both splits. Turn 2 is too eager; turn 4 wastes turns after the evidence has arrived.

**Confidence gating.** A fixed turn number is crude — sometimes the answer is obvious on turn 2. `_confident()` allows early recommendation when the top candidate clearly leads:

`(best_score - runner_up_score) / best_score >= 0.20`

| Margin         | holdout |
|----------------|---------|
| 0.0 (disabled) | 0.8349  |
| 0.15           | 0.8391  |
| 0.30           | 0.8389  |
| 0.50           | 0.8371  |

Every enabled value beats disabled, on both splits, and **MRR is unchanged while MTTC improves** — exactly the intended mechanism. The agent gets faster without ranking worse.

**The default is 0.20, not 0.15.** 0.15 scored highest, but the curve from 0.15 to 0.50 is flat, which means the differences between them are noise. Picking a split’s argmax is how you buy noise. Mid-plateau is the defensible choice — the same reasoning applies to `broad_yield` in `src/policy.py`, whose dev results (0.8530 at 5.0, 0.8651 at 8.0, 0.8634 at 12.0) are one plateau.

**Measured effect** 0.8543 → 0.8592, with MTTC improving from 3.58 to 3.41.

**Ideas for this stage**

- **A calibrated stopping rule.** The margin test is a heuristic on raw score gaps. Converting the reranker’s output into an actual probability that the top candidate is correct — fitted on the public sessions — would

let the agent stop on a meaningful threshold (“80% confident”) instead of an arbitrary ratio.

- **Spend the remaining turns.** The agent stops at the first hit, but it does not know that a hit occurred. On sessions heading for a miss it still shows the same list every turn from turn 3 onward. Detecting stagnation — the shortlist stopped changing, the customer stopped disclosing — and switching to a different strategy would attack the 6% of sessions that currently miss entirely.
- **Diversify a low-confidence list.** When confidence is low, the top 10 are often near-identical products. Deliberately spreading the list across distinct categories or brands would raise the chance of catching the target, at some cost to rank. Given that a miss→hit is worth 2.6x a rank improvement (§3), that trade may well be positive — and it is untested.
- **Vary list length with confidence.** `list_size_ramp` exists in `AgentConfig` but ships flat at 10. Since wrong recommendations are free (consequence #2), there is likely no benefit to showing fewer — but the mechanism is there and unexercised.

---

## S9 — Robustness and tests

**Files:** `starter/agent.py`, `tests/`

**What this stage does** Ensures the agent never crashes, always emits a valid response, and that changes cannot silently break the score.

**Why it was needed** Consequence #4 in §3: the evaluator swallows exceptions and turns them into missed sessions with no error shown. A crash on turn 1 of every session would produce a score of zero and an entirely clean-looking run.

**What changed** `respond()` cannot raise. The whole path is wrapped; any exception degrades that turn to a valid empty response rather than losing the session. Each stage also degrades locally — a malformed search expression returns `[]`, an illegal attribute is coerced to `other`, a missing session is recreated rather than raising.

**25 tests, in three groups:**

| File                                  | Covers                                                                                                                                                                                                   |
|---------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>tests/test_contract.py</code>   | Every response validates against <code>docs/agent_api_contract.json</code> ; hostile input (empty, 5000 characters, SQL) never raises; sessions don't leak into each other; <code>top_k</code> respected |
| <code>tests/test_components.py</code> | Each stage in isolation — including a test asserting the gain-ratio fix, so the brand-cardinality bug cannot silently return                                                                             |
| <code>tests/test_regression.py</code> | End-to-end score on the <b>holdout</b> split must stay above <b>0.78</b> — the committed floor                                                                                                           |

The regression test is the important one. The pipeline has many interacting knobs and it is easy to make a local improvement that costs score overall; this is what catches that.

```
python3 -m unittest discover -s tests -t . # 25 tests, ~14s
```

**Measured effect** No score gain. It is insurance — and it caught two real bugs during development, including a policy that kept re-asking a question the customer had already declined.

#### Ideas for this stage

- **Assert the fallbacks never fire.** Defensive layers can mask real bugs. Counting fallback activations and asserting the count is zero on a clean public run would distinguish “robust” from “quietly broken”.
- **Property-based testing.** Generating random conversations and asserting invariants (never crashes, never recommends a non-catalog ID, never re-asks a dead attribute) would cover far more input shapes than 25 hand-written cases.
- **A performance regression test.** Latency is currently measured ad hoc. Pinning per-turn latency in a test would catch an accidentally quadratic change before judging day.

## 6. What was tried and rejected

These are first-class results. They are the evidence that the shipped design was **chosen** rather than assumed, and each one is recorded in the relevant module's docstring so the next reader does not repeat the experiment.

| Experiment                                              | Result                                         | Outcome                                                                                                                                          |
|---------------------------------------------------------|------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Verbatim matching as a retrieval route</b>           | 0.6859 vs floor 0.7799                         | Removed from retrieval, moved to reranking — where it became the second-largest gain. Found the target in only 47/80 sessions as a search route. |
| <b>Routing pool size and timing by intent</b>           | dev unchanged, holdout $-0.002$                | Router scoped down to phrasing only.                                                                                                             |
| <b>Weighting phrases by rarity within the shortlist</b> | $+0.0002$ dev, 0.0000 holdout                  | Deleted entirely, not left as an off-by-default flag.                                                                                            |
| <b>InfoGainPolicy as the default</b>                    | dev 0.8369 vs 0.8738; holdout 0.8141 vs 0.8374 | Kept in the repo and documented, but not the default.                                                                                            |
| <b>Dense-vector retrieval</b>                           | not built                                      | Cancelled after measuring retrieval recall at 80/80 — the bottleneck was ranking, and it would have broken the no-network guarantee.             |

**A bug worth learning from.** The first version of `InfoGainPolicy` opened every conversation by asking the customer’s preferred *brand*. The cause was using raw entropy: `brand` has thousands of distinct values in the catalog, so it scored 6.1 bits against `colour`’s 2.3 and always won — even though it is close to useless as a shopping question. This is a well-known failure of information-gain methods, and the fix is standard: divide by the maximum possible entropy to get a **gain ratio**. `tests/test_components.py` now asserts brand scores below `colour`, so the bug cannot return.

## 7. Results

How the score was built up

| Stage of work     | Hit@10 | MRR   | MTTC | Score         |
|-------------------|--------|-------|------|---------------|
| Supplied baseline | 0.125  | 0.068 | 9.81 | <b>0.1067</b> |



| Stage of work                                              | Hit@10 | MRR   | MTTC | Score         |
|------------------------------------------------------------|--------|-------|------|---------------|
| + dialog<br>state and<br>asking<br>questions<br>(§S3, §S4) | 0.870  | 0.694 | 4.17 | <b>0.7799</b> |
| + verbatim<br>span<br>reranking<br>(§S6)                   | 0.940  | 0.786 | 3.58 | <b>0.8543</b> |
| + confidence<br>gating and<br>policy fixes<br>(§S7)        | 0.940  | 0.791 | 3.41 | <b>0.8592</b> |

(An early prototype of the state-and-questions step measured 0.7811 before the code was reorganised into modules; 0.7799 is the same policy re-measured in the final structure. The small difference comes from boilerplate word handling added in §S1.)

#### Dev versus holdout

| Split               | Sessions | Hit@10 | MRR   | MTTC | Efficiency | Score         |
|---------------------|----------|--------|-------|------|------------|---------------|
| Public (all)        | 200      | 0.940  | 0.791 | 3.41 | 0.760      | <b>0.8592</b> |
| Dev (tuned on)      | 120      | 0.950  | 0.814 | 3.27 | 0.772      | 0.8738        |
| Holdout (untouched) | 80       | 0.925  | 0.756 | 3.60 | 0.740      | 0.8374        |

Dev is higher than holdout, which is expected and healthy — dev is where the tuning happened. The gap of 0.036 is within what 80 sessions can produce by chance. The holdout number is the honest one to quote, and the one the regression test guards.

#### By scenario

| Scenario               | Sessions  | Hit@10       | MRR          | MTTC        |
|------------------------|-----------|--------------|--------------|-------------|
| boundary               | 10        | 1.000        | 0.846        | 4.00        |
| browsing               | 80        | 0.963        | 0.811        | 3.08        |
| buying                 | 80        | 0.938        | 0.794        | 3.21        |
| <b>intent_override</b> | <b>30</b> | <b>0.867</b> | <b>0.713</b> | <b>4.60</b> |

**Intent override is the weakest scenario** and the clearest target for further work. Part of the gap is structural — those sessions cannot convert before the

override arrives on turn 3 or 4 (consequence #5 in §3), which puts a floor under MTTC. The rest is genuine, and §S3’s ideas are aimed at it.

## Operational

|                             |                                                  |
|-----------------------------|--------------------------------------------------|
| Model / API                 | <b>None.</b> No LLM, no network, no credentials. |
| Dependencies                | Python standard library only                     |
| Cost                        | \$0.00                                           |
| Tokens                      | 0 prompt, 0 completion                           |
| Index build                 | 4.13s, once at startup                           |
| Per-turn latency            | 16.6ms mean, 19.8ms p95, 21.0ms max              |
| Peak memory                 | ~282 MB                                          |
| Full 200-session evaluation | ~28s                                             |

The agent behaves identically with the network disabled, because there is no online path to fall back from. This was deliberate: the submission rules reserve the right to score under network restrictions, and an agent scoring zero in that environment is worth less than one scoring 0.8592 everywhere.

## 8. Why upstream Amazon data was ruled out

The catalog derives from the public Amazon Reviews 2023 dataset, which contains far more than the 50,000 frozen products. Using it was considered and **deliberately rejected**. The reasoning, in order of strength:

- 1. It cannot add information to the channel that matters.** This is decisive. Every constraint the simulated customer can ever utter is built by `intent_card()` (`evaluator/local_evaluator.py:52`) from `features`, `details`, `price`, and `searchable_text()` over `title / features / details / description / categories / store` (`evaluator/local_evaluator.py:22`). **All of those fields are already complete in `data/catalog.jsonl`.** The process generating the dialogue is closed over the frozen catalog, so upstream data cannot tell us more about *what the customer will say* — and interpreting what they say is the entire task.
- 2. It does not address the measured bottleneck.** Retrieval already recalls the target 80/80 at median rank 1 (§S5). The losses are in ranking near-identical candidates. More product text does not break those ties.
- 3. It could actively hurt.** The traced belt failure was *too many candidates matching generic phrases*. Richer descriptions give non-target products more surface on which to match. And we have direct evidence of this class of signal backfiring: the popularity prior had to be held to weight 0.02 (§S6) because the target is one specific purchase, not a bestseller.

**4. Rule risk with a bad asymmetry.** docs/competition\_specification.md:13 places “private-label reconstruction” and “identifiers outside the frozen catalog” out of scope, and the organizer deliberately stripped review text, user IDs, timestamps and purchase history for anonymisation. Re-attaching that data by `parent_asin` works against stated intent even where it is not literally banned — risking a 0.86 result for a marginal gain.

**5. Disproportionate cost.** Tens of gigabytes of download plus a distillation script, producing an asset that must then be committed to preserve the no-network guarantee. It breaks the one-command reproduction story for very little.

**The honest hole in this argument:** `materialize_hidden_fields()` (evaluator/local\_evaluator.py:204) short-circuits when a session already carries an `intent_card`, so the private sessions *could* in principle ship cards built from richer data than the frozen catalog. That seems unlikely given the organizer froze the catalog as the shared artifact, but it cannot be verified from here.

**The legitimate half, kept.** The one genuinely useful upstream idea — mining real attribute vocabularies instead of hand-writing keyword lists — can be done against the frozen 50,000-product catalog itself. Same benefit, no download, no rule risk. It appears in §10.

---

## 9. Files at a glance

| File             | Stage           | What to look at                                                                      |
|------------------|-----------------|--------------------------------------------------------------------------------------|
| starter/agent.py | adapter, S7, S9 | <code>respond()</code> for the whole turn flow; <code>_confident()</code> for timing |
| src/index.py     | S1              | <code>CatalogIndex._build()</code> ; <code>DEFAULT_WEIGHTS</code> at line 25         |
| src/text.py      | S1              | <code>constraint_spans()</code> — how customer phrases are normalised                |
| src/router.py    | S2              | <code>classify()</code> ; the docstring records the scope-reduction measurement      |
| src/state.py     | S3              | <code>DialogState.observe()</code> ; <code>apply_override()</code> for the deviation |

| File             | Stage  | What to look at                                                                  |
|------------------|--------|----------------------------------------------------------------------------------|
| src/policy.py    | S4     | <code>FixedPolicy</code> (line 40, ships); <code>InfoGainPolicy</code> (line 72) |
| src/retrieval.py | S5     | <code>retrieve()</code> and <code>_rrf()</code>                                  |
| src/rerank.py    | S6     | <code>rerank()</code> — the span-coverage loop                                   |
| src/facets.py    | shared | <code>extract()</code> ; currently unused by ranking (see §S6 ideas)             |
| tools/sweep.py   | S0     | <code>build_configs()</code> — add a row to test a change                        |
| tests/           | S9     | <code>test_regression.py</code> guards the 0.78 floor                            |
| ARCHITECTURE.md  | —      | Data-flow diagram and stage boundaries                                           |
| README.md        | —      | Results and decision rationale                                                   |

**Never modify** `evaluator/`, `data/`, or `docs/` — they are organizer-owned and scoring is invalid if they change.

## 10. Cross-cutting and orchestration ideas

These are not tied to one stage — they concern how the stages work *together*, or they touch several at once.

### Orchestration

**Close the loop between asking and ranking.** Right now the pipeline is a one-way pass: state → retrieve → rerank → ask. The clarification policy reads the candidate pool, but nothing feeds *ranking confidence* back into *what to ask*. A genuinely adaptive agent would notice “my top two candidates differ only in colour” and ask about colour specifically — a question chosen to break the current tie rather than to split the pool in general. This is the most interesting unbuilt idea in the project, it needs no new data, and it connects S4, S6 and S7 into an actual feedback loop rather than a pipeline.

**Adaptive orchestration by session state.** Every session runs the identical sequence of stages. A session where the customer is disclosing freely and a session where they keep saying “no preference” call for different behaviour — the first should keep asking, the second should stop asking and start showing lists.

The signals to detect this already exist (`productive_turns`, `dead_attributes` in `src/state.py`); nothing consumes them for orchestration.

**A per-turn budget.** All stages run every turn regardless of whether they can contribute. The focused retrieval route only matters after an override; reranking only matters once spans exist. Skipping stages that cannot help would cut latency, and more importantly would make the agent’s reasoning legible: “I skipped reranking because nothing has been disclosed yet.”

**Explain the recommendation.** The reranker knows exactly why each candidate scored well — which phrases matched. Surfacing that in the customer-facing `message` (“these all have a stainless steel band and a date window”) costs nothing, is listed under Innovation Directions in `docs/competition_specification.md:87`, and makes the demo far more compelling than a bare list.

## Cross-cutting

**Learn the weights instead of setting them.** At least eight constants across the pipeline were set by hand or inherited: `DEFAULT_WEIGHTS` (S1), `weight_focused` (S5), `PRE_OVERRIDE_WEIGHT` (S3), `span_weight / retrieval_weight / popularity_weight / length_bonus` (S6), `confidence_margin` (S7). There are 176 known-correct public sessions to learn from. A single offline fitting pass — plain logistic regression, entirely within the standard-library and no-network constraints — would likely beat hand-tuning across the board and would be more defensible than a series of individual sweeps. **This is the highest-expected-value cross-cutting change.**

**Mine facet vocabularies from the frozen catalog.** `src/facets.py` uses hand-written keyword lists for material, colour, style, size and use case. Harvesting these from the 50,000 products themselves would be broader, less arbitrary, and would improve both the clarification policy (S4) and the proposed facet-agreement ranking signal (S6). This is the legitimate half of the rejected upstream idea, with none of its downsides.

**An error-analysis tool.** The remaining 6% of sessions that miss have never been examined as a group. A script that dumps every missed session with its disclosed constraints, the target’s actual rank, and what outranked it would likely reveal one or two systematic causes worth more than any parameter sweep. The belt case in §S6 was found this way, by hand, and immediately suggested category anchoring.

**Make robustness measurable, not assumed.** The no-network guarantee, the never-raises guarantee and the latency budget are all currently claims backed by design rather than by continuous verification. Tests that actively enforce them (§S9 ideas) would turn three assertions in this document into three facts.

## Glossary

| Term                   | One-line meaning                                                                                     |
|------------------------|------------------------------------------------------------------------------------------------------|
| <b>BM25</b>            | Text-search scoring formula. Rare words count more; matches in short important fields count more.    |
| <b>Constraint span</b> | A short phrase the customer said, normalised so it can be matched against product text.              |
| <b>Dev / holdout</b>   | 120 sessions used for tuning / 80 kept untouched to check the tuning generalises.                    |
| <b>Entropy</b>         | A measure of how evenly a set is split. High entropy = very uncertain = a useful question to ask.    |
| <b>Facet</b>           | A structured product attribute with limited values — material, colour, price band.                   |
| <b>FTS5</b>            | SQLite’s built-in full-text search engine. Requires no installation.                                 |
| <b>Gain ratio</b>      | Entropy divided by its maximum, to stop attributes with many values from always winning.             |
| <b>Hit Rate@10</b>     | Fraction of sessions where the target appeared in a shown top-10 list.                               |
| <b>IDF</b>             | Inverse document frequency — the “rare words matter more” weighting inside BM25.                     |
| <b>MRR</b>             | Mean Reciprocal Rank — average of $1/\text{position}$ . Rewards ranking first, not merely appearing. |
| <b>MTTC</b>            | Mean Turns To Conversion — average turn on which the hit happened; a miss counts as 11.              |
| <b>Overfitting</b>     | Tuning until you fit the noise in your test data, so results don’t transfer to new data.             |
| <b>parent_asin</b>     | The product ID. The only thing scored, and only on exact match.                                      |
| <b>Provenance</b>      | Metadata about <i>when and under what circumstances</i> a piece of information arrived.              |
| <b>Reciprocal rank</b> | $1 / \text{position}$ . Rank 1 $\rightarrow$ 1.0, rank 10 $\rightarrow$ 0.1, miss $\rightarrow$ 0.   |

| Term             | One-line meaning                                                                    |
|------------------|-------------------------------------------------------------------------------------|
| <b>Reranking</b> | Reordering an existing shortlist to put the best candidate first.                   |
| <b>Retrieval</b> | Finding a shortlist of candidates from the full catalog.                            |
| <b>RRF</b>       | Reciprocal Rank Fusion — combining ranked lists using positions rather than scores. |
| <b>Stopword</b>  | A common word removed before searching because it carries no signal.                |
| <b>Target</b>    | The hidden product the agent must find.                                             |
| <b>Top-k</b>     | The first $k$ items of a ranked list; here $k = 10$ .                               |